

Introduction

In []:

```
'''  
Weekend Rainfall Analysis in Buenos Aires (2020-2025)  
  
This notebook investigates a common belief:  
"It rains more on weekends"  
  
Using daily precipitation data from the Buenos Aires Central Observatory, we aim  
1. Does daily precipitation (mm/day) differ between weekdays and weekends?  
2. Is the probability of rainfall different between the two?  
'''
```

Load and prepare data

In [481...]

```
#Carga de módulos  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from scipy.stats import mannwhitneyu  
import statsmodels.api as sm
```

In [482...]

```
'''  
Load of 'export' dataset, downloaded from meteostat.net under the following para  
Observatory: Buenos Aires Central Observatory  
Date range: 13/11/2020 - 13/11/2025  
  
Field information:  
date (YYYY-mm-dd)  
tavg: average Temperature (°C)  
tmin: minimum temperature (°C)  
tmax: maximum temperature (°C)  
prcp: total precipitation (mm)  
snow: snow depth  
wdir: wind (From) direction (°)  
wspd: wind speed (km/h)  
wpgt: wind peak gust (km/h)  
pres: sea-level air pressure (hPa)  
tsun: total sunshine duration (minutes)  
'''  
  
df = pd.read_csv('export.csv')  
df['date'] = pd.to_datetime(df['date'], origin='1899-12-30', unit='D')  
df.head()
```

Out[482...]

	date	tavg	tmin	tmax	prcp	snow	wdir	wspd	wpgt	pres	tsun
0	2020-11-13	22.2	19.2	26.6	NaN	NaN	NaN	9.0	NaN	1013.1	NaN
1	2020-11-14	22.2	19.3	26.6	NaN	NaN	NaN	4.0	NaN	1010.3	NaN
2	2020-11-15	20.9	15.7	26.0	NaN	NaN	NaN	11.1	NaN	1013.5	NaN
3	2020-11-16	20.5	14.0	26.3	NaN	NaN	NaN	6.8	NaN	1014.6	NaN
4	2020-11-17	23.4	17.7	27.7	NaN	NaN	NaN	8.2	NaN	1013.0	NaN

In [483...]

```
print('\nStatistical description of data:\n')
print(df.describe(include='all'))
```

Statistical description of data:

	date	tavg	tmin	tmax	\		
count	1827	1827.000000	1827.000000	1827.000000			
mean	2023-05-14 23:59:59.999999744	18.746196	14.418993	24.374330			
min	2020-11-13 00:00:00	3.700000	-1.900000	9.200000			
25%	2022-02-12 12:00:00	14.100000	9.800000	19.400000			
50%	2023-05-15 00:00:00	18.900000	14.600000	24.400000			
75%	2024-08-13 12:00:00	23.400000	19.100000	29.400000			
max	2025-11-13 00:00:00	33.500000	29.800000	41.500000			
std	NaN	5.840665	5.899633	6.260565			
	prcp	snow	wdir	wspd	wpgt	pres	tsun
count	1318.000000	5.000000	0.0	1827.000000	0.0	1827.000000	0.0
mean	3.986419	2.800000	NaN	8.743733	NaN	1015.985495	NaN
min	0.000000	1.000000	NaN	1.100000	NaN	997.400000	NaN
25%	0.000000	2.000000	NaN	6.400000	NaN	1011.800000	NaN
50%	0.000000	2.000000	NaN	8.400000	NaN	1015.700000	NaN
75%	1.600000	2.000000	NaN	10.700000	NaN	1019.850000	NaN
max	127.000000	7.000000	NaN	25.000000	NaN	1034.800000	NaN
std	11.128690	2.387467	NaN	3.281035	NaN	6.091385	NaN

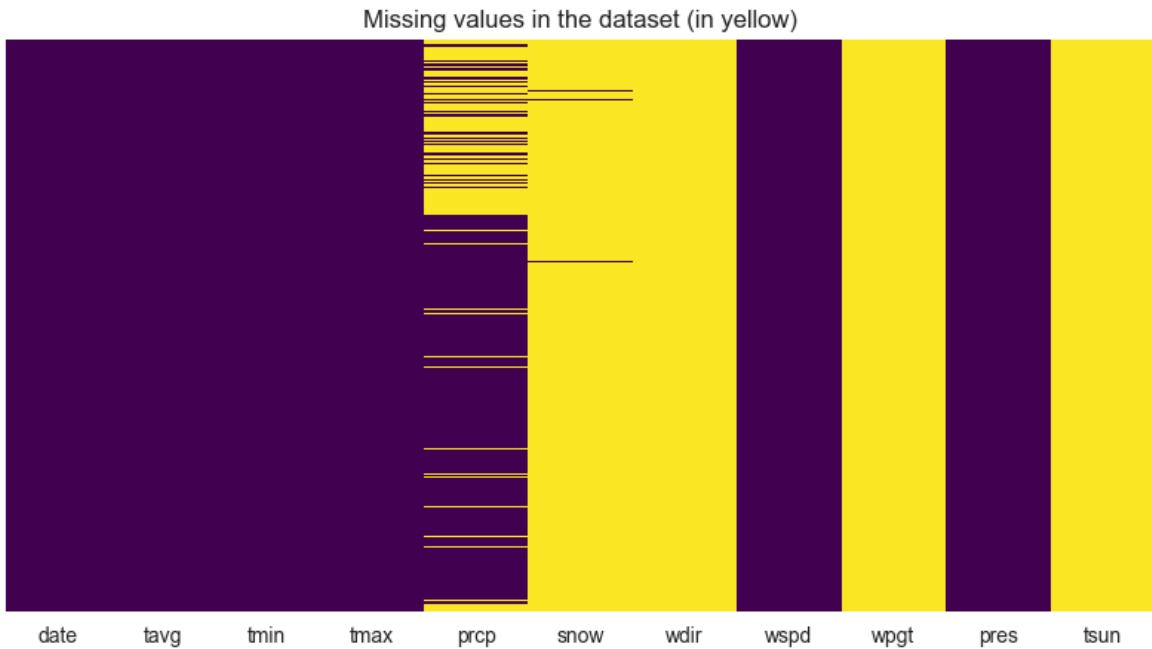
In [484...]

```
print('\nMissing values by field:\n')
print(df.isna().sum())

plt.figure(figsize=(10,5))
sns.heatmap(df.isna(), cmap='viridis', cbar=False, yticklabels=False)
plt.title('Missing values in the dataset (in yellow)')
plt.show()
```

Missing values by field:

date	0
tavg	0
tmin	0
tmax	0
prcp	509
snow	1822
wdir	1827
wspd	0
wpgt	1827
pres	0
tsun	1827
dtype: int64	



Dataset cleaning

In [485...]

```
...
Are there fields that should be eliminated?
-wdir, wpgt and tsun have no non-null values.
-in wspd, snow and pres, the non-null values represent more than 60% of the total
...
cols_to_drop=['snow', 'wdir', 'wspd', 'wpgt', 'pres', 'tsun']
df = df.drop(columns=[c for c in cols_to_drop if c in df.columns])

df.head()
```

Out[485...]

	date	tavg	tmin	tmax	prcp
0	2020-11-13	22.2	19.2	26.6	NaN
1	2020-11-14	22.2	19.3	26.6	NaN
2	2020-11-15	20.9	15.7	26.0	NaN
3	2020-11-16	20.5	14.0	26.3	NaN
4	2020-11-17	23.4	17.7	27.7	NaN

Null value handling

In [486...]

```
...
Upon reaching our main field of interest (precipitation), we see that there is a
Are they mostly corresponding to days without rain or they have missing data for
To determine it, the daily precipitation is grouped by year (2021-2024) and cont
...
#Dataset preparation with daily records. Creation of columns 'year' and 'rainy_d
```

```
df['year'] = df['date'].dt.year
df['rainy_day'] = df['prcp']>0

df_review = df.groupby('year').agg(
```

```

        daily_mm = ('prcp', 'sum'),
        daily_rainy_days= ('rainy_day', 'sum'),
    )

#Loading the OCBA dataset with annual data.
ocba = pd.read_csv('ocba.csv')
ocba.rename(columns={'mm':'ocba_mm', 'days':'ocba_rainy_days'}, inplace=True)

#Inner join of both datasets.

df_merge = df_review.merge(ocba, left_on='year', right_on='year', how='inner')
...
To continue, the relative error of the dataset is calculated, taking as the sour
An acceptable relative error of up to +/-5% is considered. Starting from that va
...
df_merge['prcp_difference']=((df_merge['daily_mm']-df_merge['ocba_mm'])/df_merge
df_merge['rainy_days_difference']=((df_merge['daily_rainy_days']-df_merge['ocba_'

df_merge.head()

```

Out[486...]

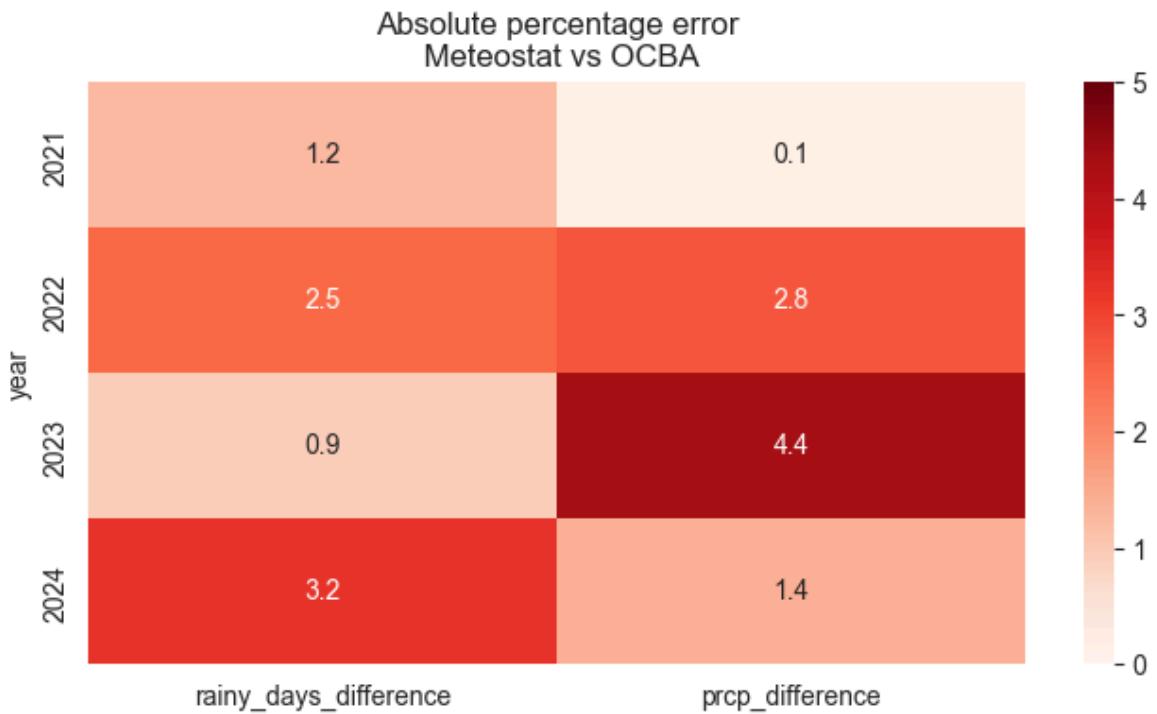
	year	daily_mm	daily_rainy_days	ocba_mm	ocba_rainy_days	prcp_difference	rainy_
0	2021	961.0		79	959.8	80	0.125026
1	2022	773.2		82	752.4	80	2.764487
2	2023	912.5		112	954.0	111	-4.350105
3	2024	1145.4		96	1161.7	93	-1.403116

In [487...]

```

plt.figure(figsize=(8,4))
df_abs = df_merge.abs()
sns.heatmap(
    df_abs.set_index('year')[['rainy_days_difference', 'prcp_difference']],
    annot=True, fmt='.1f', cmap='Reds', vmin=0,vmax=5)
plt.title('Absolute percentage error\n Meteostat vs OCBA')
plt.show()

```



```
In [488...]: #As can be observed both analytically and graphically, the absolute relative error
df['prcp'] = df['prcp'].fillna(0)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1827 entries, 0 to 1826
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   date        1827 non-null    datetime64[ns]
 1   tavg         1827 non-null    float64 
 2   tmin         1827 non-null    float64 
 3   tmax         1827 non-null    float64 
 4   prcp         1827 non-null    float64 
 5   year         1827 non-null    int32   
 6   rainy_day    1827 non-null    bool    
dtypes: bool(1), datetime64[ns](1), float64(4), int32(1)
memory usage: 80.4 KB
```

Outlier Detection

```
In [489...]: ...
The appearance of outliers can affect the final conclusions of the analysis, so
possible atypical values in both a graphical and analytical manner.
Given that the nature of a precipitation time series has an extremely high level
is not normal, is considerably asymmetric and has a long tail, only the data wit
...
sns.boxplot(x=df[df['rainy_day'] == True]['prcp'])
plt.title('Precipitation distribution for rainy days')

#Values above the 99.5 percentile
p995 = df[df['rainy_day'] == True]['prcp'].quantile(0.995)
df_outliers = df.loc[df['prcp'] > p995, ['date', 'prcp']]
df_neg = df.loc[df['prcp'] < 0, ['date', 'prcp']]
```

```
print(f'Values above the 99.5 percentile\n\n {df_outliers}\n')
print(f'Negative precipitations\n\n {df_neg}')
```

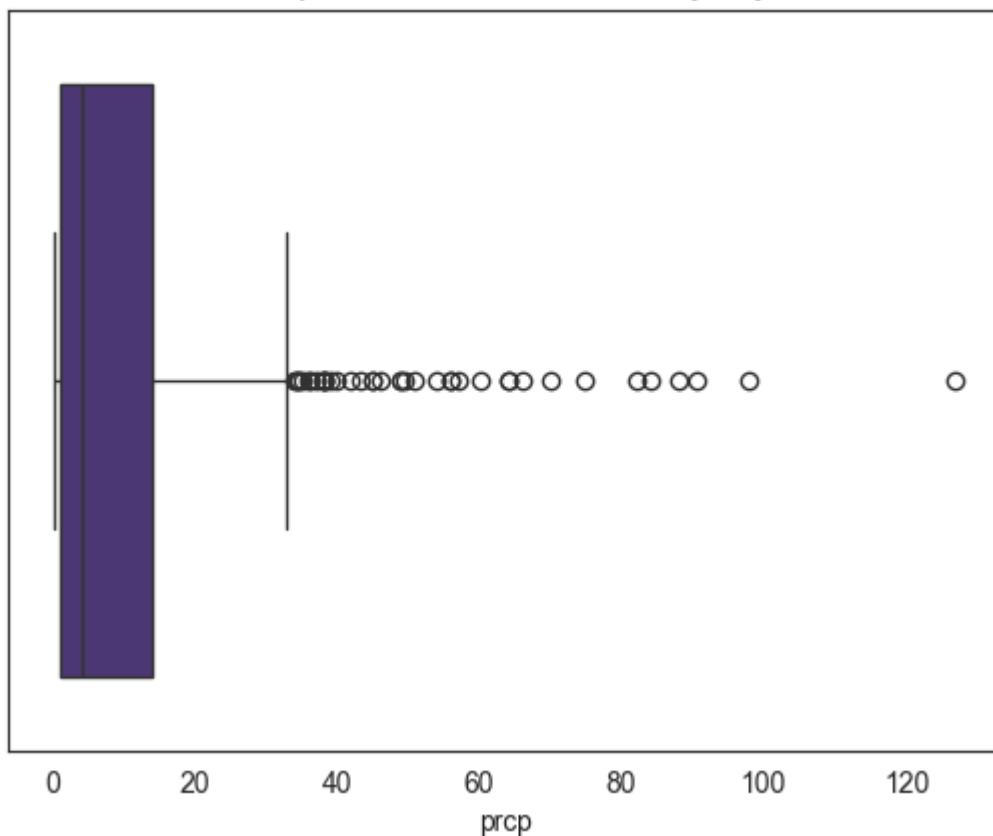
Values above the 99.5 percentile

	date	prcp
731	2022-11-14	90.5
1215	2024-03-12	127.0
1646	2025-05-17	98.0

Negative precipitations

Empty DataFrame
Columns: [date, prcp]
Index: []

Precipitation distribution for rainy days



In []:

```
...
What conclusions can be drawn from the results obtained?  
When analyzing a meteorological variable such as precipitation, typical values m  
In this case, the three detected days were compared with the SMN records on the  
  
Inevitably erroneous values would, in this case, be negative values or values gr  
In this analysis, none of these two types of errors have been found.  
...
```

Dataset transformation

In [491...]

```
...
Fields are created to discriminate records by type of day of the week:  
day_of_week indicates the weekday number, starting from index 0 (Monday)  
is_weekend is a boolean variable that returns True for Friday, Saturday, and Sun  
...
```

```

df['day_of_week'] = df['date'].dt.dayofweek
df['is_weekend'] = df['day_of_week']>3
df.head()

```

Out[491...]

	date	tavg	tmin	tmax	prcp	year	rainy_day	day_of_week	is_weekend
0	2020-11-13	22.2	19.2	26.6	0.0	2020	False	4	True
1	2020-11-14	22.2	19.3	26.6	0.0	2020	False	5	True
2	2020-11-15	20.9	15.7	26.0	0.0	2020	False	6	True
3	2020-11-16	20.5	14.0	26.3	0.0	2020	False	0	False
4	2020-11-17	23.4	17.7	27.7	0.0	2020	False	1	False

Hypothesis testing

In [492...]

```

'''
It's time to test our two questions:
1) Is there really a higher probability of rain on the weekend than on the rest
2) On average, does it rain more on the weekend than on the rest of the days?

Let's start by grouping our data and examining the results.
'''

prcp_by_weekday = df.groupby('is_weekend').agg(
    prcp = ('prcp', 'mean'),
    rainy_days = ('rainy_day', 'mean'),
    total_days = ('is_weekend', 'count')
)

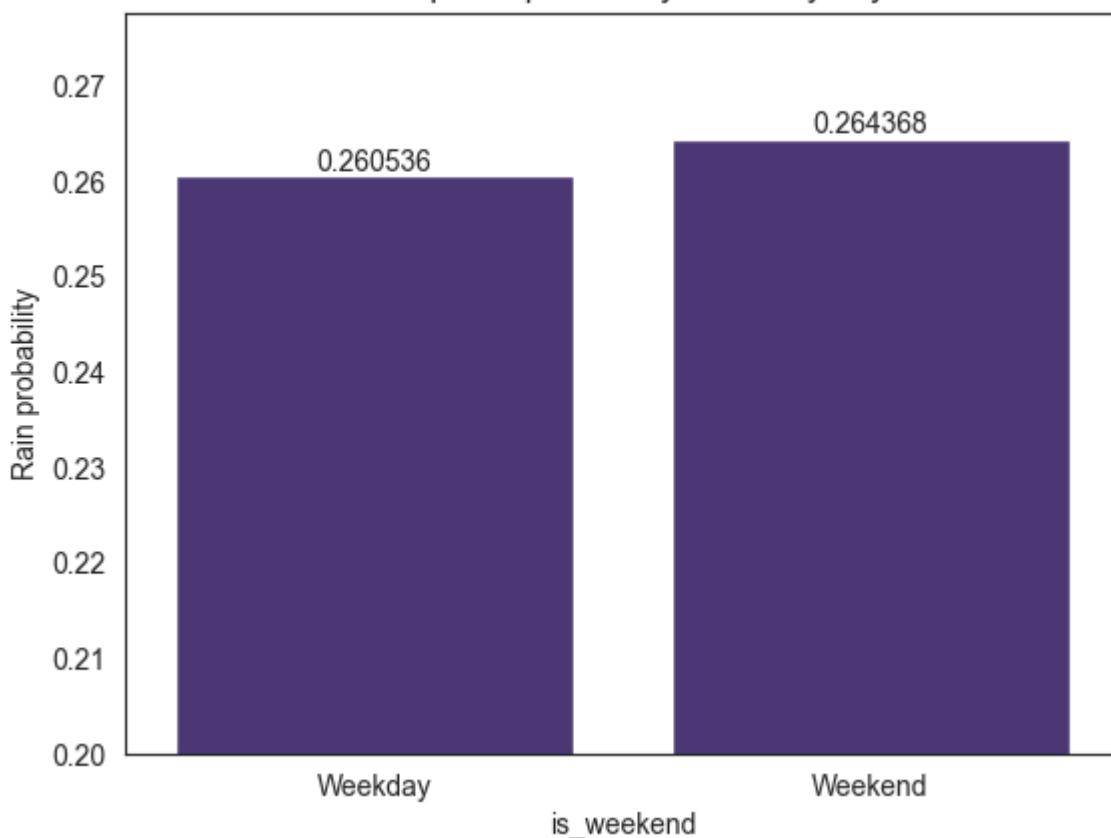
ax = sns.barplot(
    data=prcp_by_weekday.reset_index(),
    x="is_weekend",
    y="rainy_days"
)
ax.set_ylim(bottom=0.20)
ax.bar_label(ax.containers[0])
plt.xticks([0,1], ["Weekday", "Weekend"])
plt.ylabel("Rain probability")
plt.title("Compared probability of a rainy day")
plt.show()

sns.pointplot(
    data=prcp_by_weekday.reset_index(),
    x="is_weekend",
    y="prcp",
    capsize=0.2
)
plt.xticks([0,1], ["Weekday", "Weekend"])
plt.ylabel("average mm per day")
plt.title("Average rainfall (with CI)")
plt.show()

prcp_by_weekday

```

Compared probability of a rainy day



Average rainfall (with CI)



Out[492...]

prcp rainy_days total_days

is_weekend	prcp	rainy_days	total_days
False	2.706418	0.260536	1044
True	3.101660	0.264368	783

Is there really a higher probability of rain on the weekend than on the rest of the days?

In [493...]

'''
 We can observe that, in Buenos Aires, it rained on 26.44% of weekend days compared to 26.05% of weekday days. Although the percentage difference seems small, we want to determine whether this difference is statistically significant. To do so, we will perform the well-known two-proportion Z-test, which evaluates whether the proportions are significantly different. In our case, the groups are weekend (Fri-Sun) and weekdays (Mon-Thu).

Statements:

Null Hypothesis (H_0): The proportion of rainy days is the same for weekends and weekdays.
 Alternative Hypothesis (H_1): The proportion of rainy days differs between weekend and weekday.

```
'''  

# We will begin the test by creating Boolean series determining whether it is a  

rain_weekday = df[df['is_weekend'] == False]['rainy_day']  

rain_weekend = df[df['is_weekend'] == True]['rainy_day']  

count = [rain_weekend.sum(), rain_weekday.sum()] #number of successes (rain)  

nobs = [len(rain_weekend), len(rain_weekday)] #number of days in each group  

z_stat, p_value = sm.stats.proportions_ztest(count, nobs)  

print('p-value for the Z-test:', p_value)
```

p-value for the Z-test: 0.8538041244079058

In []:

'''
 The p-value is the probability of obtaining the results observed in a study, assuming the null hypothesis is true. A p-value less than or equal to 0.05 is typically considered statistically significant. In this case, a p-value of 0.8538 is much greater than 0.05. In other words, there is no evidence to reject H_0 , so we cannot conclude that there is a significant difference in the proportion of rainy days between weekends and weekdays.

Is there any difference in daily precipitation between the weekend and the rest of the days?

In [495...]

'''
 To answer this question, we will compare the precipitation distributions between weekend and weekday days. Given the nature of a precipitation distribution (previously explained), the Mann-Whitney U test is more appropriate than a t-test. This non-parametric test (does not assume normality) is indicated for data with different distributions. In this case, we state:

Null Hypothesis (H_0): There is no difference in the daily precipitation distribution between weekend and weekday days.
 Alternative Hypothesis (H_1): There is a difference in the daily precipitation distribution between weekend and weekday days.

```
'''  

#We create the x and y series to separate precipitations into the two groups  

x = df[df["is_weekend"] == False]["prcp"] #weekdays  

y = df[df["is_weekend"] == True]["prcp"] #weekends  

#Test  

u_stat, p_value = mannwhitneyu(x, y, alternative='two-sided')
```

```
print("p-value:", p_value)
```

```
p-value: 0.61690366009889
```

```
In [ ]: """
With a p-value of 0.6170, we cannot reject H0. That is, there is no evidence that
"""
```

General Conclusions

```
In [ ]: """
In this analysis, we asked whether there is truly a greater chance of rain on th
To make the study more robust, we also wanted to see whether there was a differe
To answer these questions, we followed the classic processes of exploratory data
handling of missing data, and selection of relevant variables.
```

```
Then, we formally stated our hypotheses and performed the percentile statistical
In this case, both hypotheses were raised for both cases, so the differences fou
are due to the natural variability of the climate for the period 2020-2025 in Bu
```

```
We can say that the perception that it rains more likely when the weekend arrive
than to a climatological pattern.
"""
```